
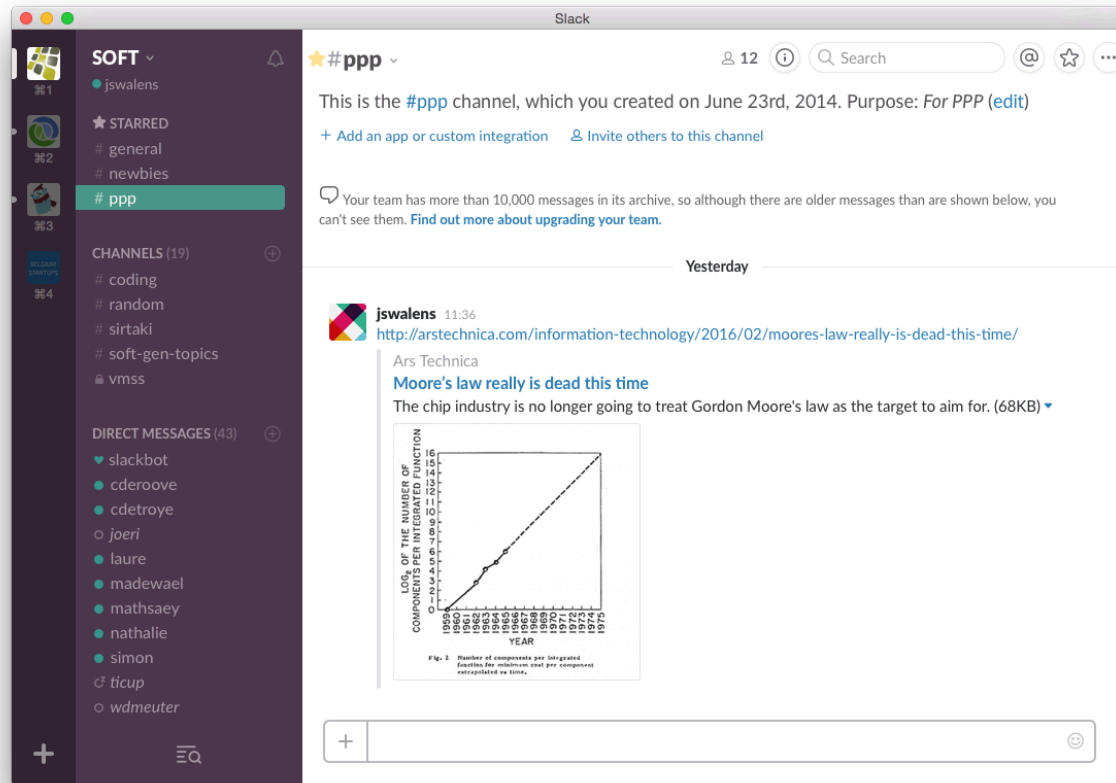


Multicore Programming
Project Erlang
Slack

Scalable chat room (Slack)



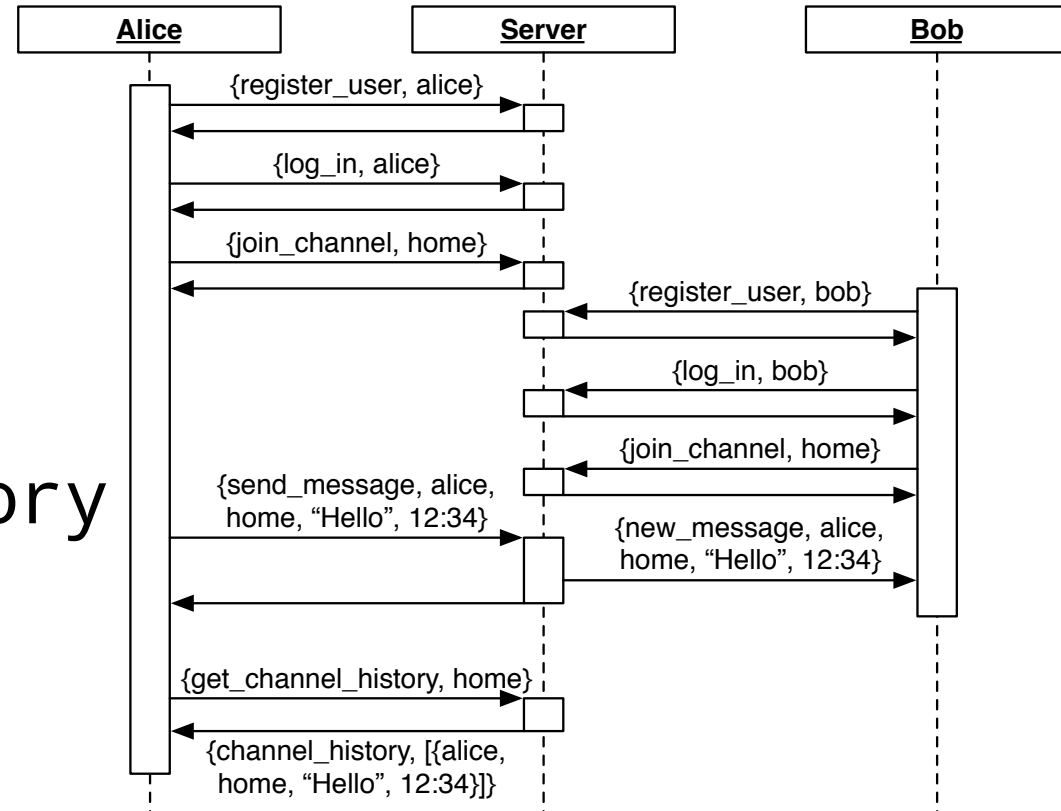
API

Client → server:


- register_user
- log_in
- join_channel
- send_message
- get_channel_history

Server → client:

- new_message



Example Implementation



- ◆ 1 process / client → don't change, for benchmarking
- ◆ 1 server process → you'll need to change this

Scalability over Consistency

Allowed to sacrifice consistency for scalability

E.g:

- ◆ Some requests might return stale data
- ◆ Data may be duplicated and be inconsistent for short periods
- ◆ Keep both channels for user as well as users for channel, both lists sometimes inconsistent

Describe motivations in report


Evaluation

Benchmarks: generate number of client processes, each do a number of operations (send_message, get_channel_history)

Latency & throughput of operations, in variety of situations

Best case, worst case, in between

Allowed to share benchmark code on PointCarré forum (only this, nothing else!)



Evaluation

Serenity = 64-core server at lab

Time slot between March 21st and April 6st
(will send out Doodle)

Remote log in: demo in lab session

Report may contain experiments on Serenity as well
as local (≥ 4 cores)

Report

Table of contents in assignment sheet:

Implementation

How do you ensure scalability?

Where did you sacrifice consistency?

Include diagrams to show design

Evaluation

Describe set-up, metrics

Include plots of results

Insight questions

Hypothetical extensions

Details

Deadline: Friday 8th of April, 23:59

Submit ZIP with implementation & report on PointCarré

Project defense in June

$\frac{1}{3}$ of your final grade

Assignment sheet on PointCarré